

Lesson 07: Sampling and Sampling Distributions

References

- Black, Chapter 7 Sampling and Sampling Distributions (pp. 201-223)
- Davies, Chapter 17 Sampling Distributions and Confidence (pp. 367-377)

Exercises:


- 1) Use the uniform distribution over the interval 0 to 1. Draw 100 random samples of size 10. Calculate the means for each sample. Using the 100 mean values plot a histogram. Repeat with 100 random samples of size 50. Repeat with 100 samples of size 500. Present the three histograms using `par()`. Calculate the variance of each histogram and compare to the original uniform distribution. What do you conclude?

```
set.seed(1234)

x <- c() # creates empty vector
y <- c() # creates empty vector

for (i in 1:100) {
  z <- runif(10)
  x <- append(x, mean(z)) # vector "x" will contain our 100 means
  y <- append(y, var(z))  # vector "y" will contain our 100 variances
}

#plot histogram of means
hist(x)
```




```

# repeat with 100 samples of size 50
a <- c() # creates empty vector
b <- c() # creates empty vector

for (i in 1:100) {
  c <- runif(50)
  a <- append(a, mean(c))
  b <- append(b, var(c))
}

#plot histogram of means
hist(a)

```




```

# repeat with 100 samples of size 50
d <- c() # creates empty vector
e <- c() # creates empty vector

for (i in 1:100) {
  f <- runif(500)
  d <- append(d, mean(f))
  e <- append(e, var(f))
}

#plot histogram of means
hist(d)


```



```
# present the three histograms using par() and mfrow()
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(x, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
hist(a, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
hist(d, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
mtext("Histogram of random uniform sample means
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random uniform sample means\n\t(n = 10, n =
## 50 and n = 500)", : font width unknown for character 0x9
```

Histogram of random uniform sample means
($n = 10$, $n = 50$ and $n = 500$)




```
par(mfrow=c(1,1))

# create and present histograms of the sample variances
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(y, main = "", xlim = c(0.02,0.16))
abline(v = 0.0833333, col = "darkred", lwd = 2)
hist(b, main = "", xlim = c(0.02,0.16))
abline(v = 0.0833333, col = "darkred", lwd = 2)
hist(e, main = "", xlim = c(0.02,0.16))
abline(v = 0.0833333, col = "darkred", lwd = 2)
mtext("Histogram of random uniform sample variances
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random uniform sample variances\n\t(n = 10,
## n = 50 and n = 500)", : font width unknown for character 0x9
```

Histogram of random uniform sample variances (n = 10, n = 50 and n = 500)



```
par(mfrow=c(1,1))
```

```
# NOTE: abline() for mean and variance histograms equal to "true" values
# for a uniform distribution (0,1).
```

```
# mean = a + b / 2 = 0.5
```

```
# variance = (b - a)^2 / 12 = 0.08333333
```

- 2) Using the histogram determined above for samples of size 50, find the quartiles. Using the normal distribution with the true mean and variance for a uniform distribution over the interval 0 to 1, determine the theoretical quartiles for a sample mean from 50 observations. Compare the two sets of quartiles. What do you conclude?

```
j <- c()
```

```
for (i in 1:100) {
  k <- rnorm(50, 0.5, sqrt(0.08333333))
  j <- append(j, mean(k))
}
```

```
quantile(a)
```

```
##          0%          25%          50%          75%          100%
## 0.4021239 0.4674217 0.5029306 0.5290173 0.5946016
```

```
quantile(j)
```

```
##          0%          25%          50%          75%          100%
## 0.4158362 0.4810670 0.5022179 0.5252440 0.6535995
```


```
# The two sets of quartiles are very similar, likely to converge as sample sizes
# are increased.
```

- 3) Use the binomial distribution with $p = 0.5$. Draw 100 random samples of size 10. Calculate the means for each sample. Using the 100 mean values plot a histogram. Repeat with 100 random samples of size 50. Repeat with 100 samples of size 500. Present the three histograms using `par()`. Calculate the variance of each histogram and compare to the original mean and variance for the binomial. What do you conclude?

```
m <- c()
n <- c()

for (i in 1:100) {
  o <- rbinom(10, 1, p = 0.5)
  m <- append(m, mean(o))
  n <- append(n, var(o))
}


# plot histogram of means
hist(m)
```



```
# repeat with 100 samples of size 50
p <- c()
q <- c()

for (i in 1:100) {
  r <- rbinom(50, 1, p = 0.5)
  p <- append(p, mean(r))
  q <- append(q, var(r))
}


hist(p)
```



```
# repeat with 100 samples of size 500
s <- c()
t <- c()

for (i in 1:100) {
  u <- rbinom(500, 1, p = 0.5)
  s <- append(s, mean(u))
  t <- append(t, var(u))
}


hist(s)
```



```
# present the three histograms using par() and mfrow()
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(m, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
hist(p, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
hist(s, main = "", xlim = c(0.2,0.8))
abline(v = 0.5, col = "darkred", lwd = 2)
mtext("Histogram of random binomial (p = 0.5) sample means
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random binomial (p = 0.5) sample means\n\t(n
## = 10, n = 50 and n = 500)", : font width unknown for character 0x9
```


Histogram of random binomial ($p = 0.5$) sample means
($n = 10$, $n = 50$ and $n = 500$)




```
par(mfrow=c(1,1))

# create and present histograms of the sample variances
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(n, main = "", xlim = c(0.10,0.25))
abline(v = 0.25, col = "darkred", lwd = 2)
hist(q, main = "", xlim = c(0.10,0.25))
abline(v = 0.25, col = "darkred", lwd = 2)
hist(t, main = "", xlim = c(0.10,0.25))
abline(v = 0.25, col = "darkred", lwd = 2)
mtext("Histogram of random binomial (p = 0.5) sample variances
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random binomial (p = 0.5) sample variances\n
## \t(n = 10, n = 50 and n = 500)", : font width unknown for character 0x9
```

Histogram of random binomial ($p = 0.5$) sample variances ($n = 10$, $n = 50$ and $n = 500$)



```
par(mfrow=c(1,1))
```

- 4) Using the histogram determined above for samples of size 50, find the quartiles. Using the normal distribution with the true mean and variance for a binomial distribution with $p = 0.5$, determine the theoretical quartiles for a sample mean from 50 observations. Compare the two sets of quartiles. What do you conclude?

```
v <- c()

for (i in 1:100) {
  w <- rnorm(50, 0.5, sqrt(0.25))
  v <- append(v, mean(w))
}
```

```
quantile(p)
```

```
##    0%  25%  50%  75% 100%
## 0.32 0.46 0.52 0.56 0.66
```

```
quantile(v)
```

```
##          0%          25%          50%          75%          100%
## 0.3328496 0.4511897 0.4955020 0.5505300 0.6580152
```

```
# The two sets of quartiles are very similar, likely to converge as sample sizes
# are increased.
```

- 5) Use the binomial distribution with $p = 0.1$. Draw 100 random samples of size 10. Calculate the means for each sample. Using the 100 mean values plot a histogram. Repeat with 100 random samples of size 50. Repeat with 100 samples of size 500. Present the three histograms using `par()`. Calculate the variance of each histogram and compare to the original mean and variance for the binomial. What do you conclude?


```

mm <- c()
nn <- c()

for (i in 1:100) {
  oo <- rbinom(10, 1, p = 0.1)
  mm <- append(mm, mean(oo))
  nn <- append(nn, var(oo))
}

# plot histogram of means
hist(mm)

```




```

# repeat with 100 samples of size 50
pp <- c()
qq <- c()

for (i in 1:100) {
  rr <- rbinom(50, 1, p = 0.1)
  pp <- append(pp, mean(rr))
  qq <- append(qq, var(rr))
}

hist(pp)


```



```
# repeat with 100 samples of size 500
ss <- c()
tt <- c()

for (i in 1:100) {
  uu <- rbinom(500, 1, p = 0.1)
  ss <- append(ss, mean(uu))
  tt <- append(tt, var(uu))
}


hist(ss)
```



```
# present the three histograms using par() and mfrow()
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(mm, main = "", xlim = c(0,0.5))
abline(v = 0.1, col = "darkred", lwd = 2)
hist(pp, main = "", xlim = c(0,0.5))
abline(v = 0.1, col = "darkred", lwd = 2)
hist(ss, main = "", xlim = c(0,0.5))
abline(v = 0.1, col = "darkred", lwd = 2)
mtext("Histogram of random binomial (p = 0.1) sample means
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random binomial (p = 0.1) sample means\n\t(n
## = 10, n = 50 and n = 500)", : font width unknown for character 0x9
```

Histogram of random binomial ($p = 0.1$) sample means
($n = 10$, $n = 50$ and $n = 500$)




```
par(mfrow=c(1,1))

# create and present histograms of the sample variances
par(mfrow=c(1,3), oma=c(0,0,2,0))
hist(nn, main = "", xlim = c(0,0.30))
abline(v = 0.09, col = "darkred", lwd = 2)
hist(qq, main = "", xlim = c(0,0.30))
abline(v = 0.09, col = "darkred", lwd = 2)
hist(tt, main = "", xlim = c(0,0.30))
abline(v = 0.09, col = "darkred", lwd = 2)
mtext("Histogram of random binomial (p = 0.1) sample variances
      (n = 10, n = 50 and n = 500)",side = 3, outer = T, line = -1)
```

```
## Warning in mtext("Histogram of random binomial (p = 0.1) sample variances\n
## \t(n = 10, n = 50 and n = 500)", : font width unknown for character 0x9
```

Histogram of random binomial ($p = 0.1$) sample variances
($n = 10$, $n = 50$ and $n = 500$)



```
par(mfrow=c(1,1))
```

- 6) Using the histogram determined above for samples of size 50, find the quartiles. Using the normal distribution with the true mean and variance for a binomial distribution with $p = 0.1$, determine the theoretical quartiles for a sample mean from 50 observations. Compare the two sets of quartiles. What do you conclude?

```
vv <- c()

for (i in 1:100) {
  ww <- rnorm(50, 0.1, sqrt(0.09))
  vv <- append(vv, mean(ww))
}
```

```
quantile(pp)
```

```
##      0%   25%   50%   75%  100%
## 0.02 0.08 0.12 0.14 0.26
```

```
quantile(vv)
```

```
##           0%           25%           50%           75%           100%
## -0.02883012  0.08130462  0.10593952  0.13540200  0.19509431
```

```
# The two sets of quartiles are very similar, likely to converge as sample sizes
# are increased.
```